

# Introducing **lakeFS**

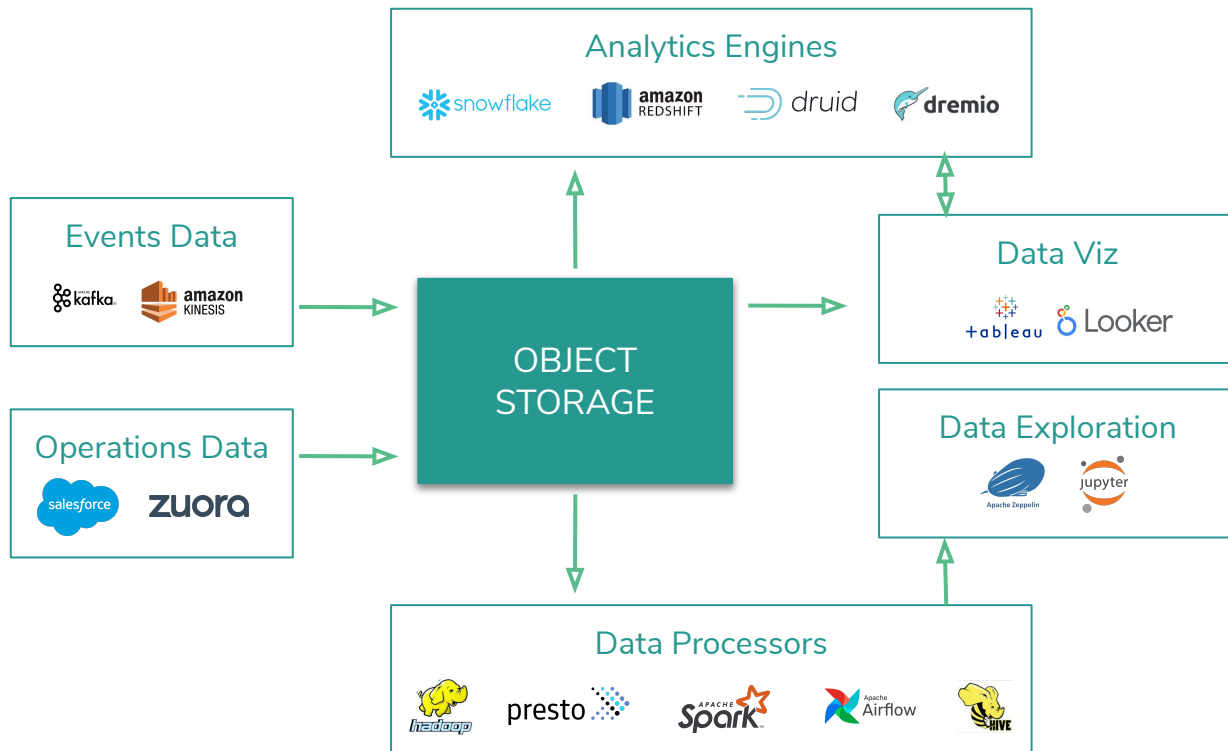
---

April 2021



**lakeFS**

# Data Lake Architecture Reference



# The Data Lake **Advantage**

1. **Scalability and cost effectiveness**
2. **Accessibility and ease of use**
3. **High throughput**
4. **Rich application ecosystem**



# The Data Lake Challenge

1. **In-ability to experiment, compare and reproduce**

Example: add new metric for BI

2. **Difficult to enforce data best practices**

Example: schema, format enforcement

3. **Hard to ensure high quality data**

Example: validate statistical properties of the data



In a perfect world

we would manage data  
from dev to production the  
way we manage code





Atomic versioned  
data lake on top of  
object storage



# Manageability & Resilience Layer

## Data Sources



## Object Store



## Data Consumption




# How it works


**API Compatibility**  
with object stores



**Versioning Engine**  
Transactional, atomic,  
isolated



**Git Terminology**  
Branch, commit, merge





# How it works

s3://data-bucket/collections/foo

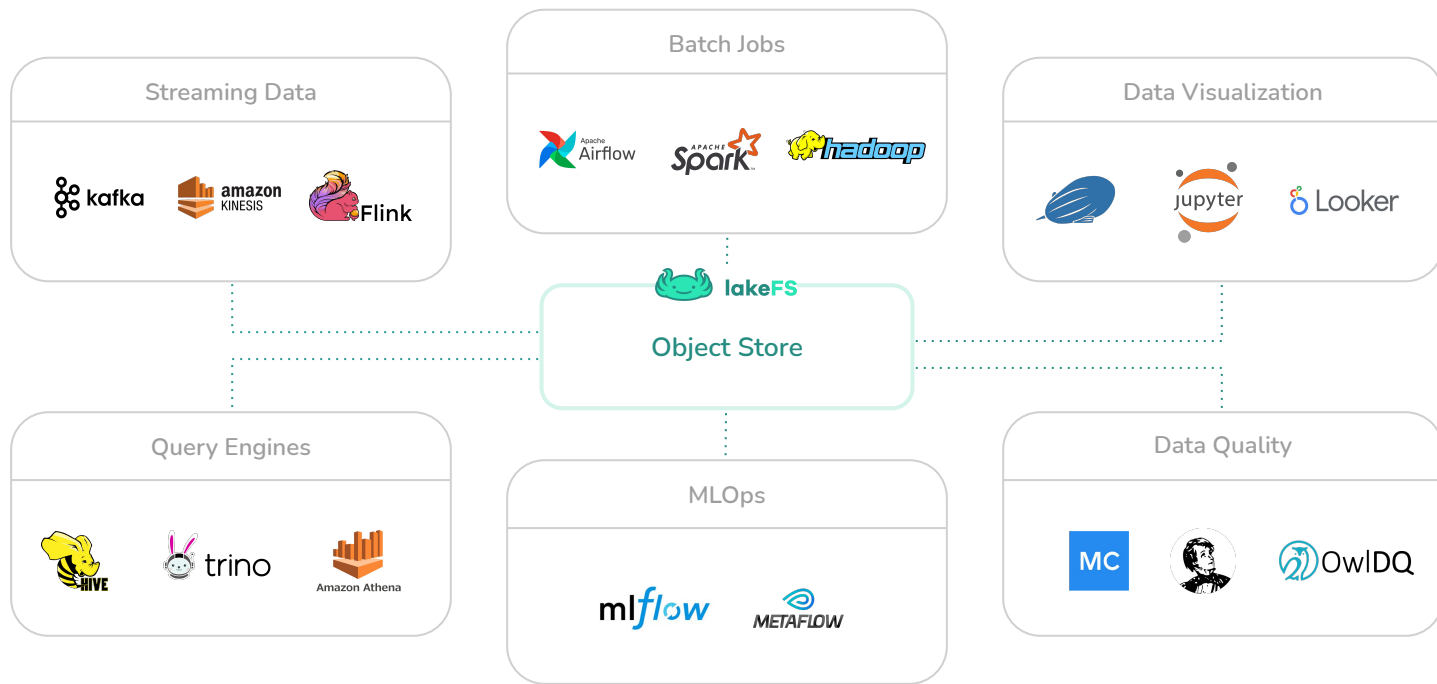


s3://data-bucket/main/collections/foo

```
# Will create an external table using the same partitions and configurations  
# but pointing at paths under the new branch  
Lakectl metastore copy \  
  --from-table user_events \  
  --to-branch dev-experiment-1
```



# Integrates with your existing tools




# Development Environment for Data

- **Experimentation** - try tools and code in isolation
- **Reproducibility** - go back to any point of time for both your code and your data
- **Compare** - tools, code or different versions of your data



# Experiment Safely



# Experimenting with Spark

```
lakectl branch create \  
  lakefs://example-repo@testing-spark-3 \  
  --source lakefs://example-repo@main  
# output:  
# created branch 'testing-spark-3.0', pointing to commit ID: 'd1e9adc71c10a'
```

```
val dfExperiment1 = sc.read.parquet("s3a://example-repo/experiment-1/events/by-date")  
val dfExperiment2 = sc.read.parquet("s3a://example-repo/experiment-2/events/by-date")  
  
dfExperiment1.groupBy(" ... ").count()  
dfExperiment2.groupBy(" ... ").count() // now we can compare the properties of the data itself
```



# Experimenting with Presto

```
lakectl branch create \  
  lakefs://example-repo@testing-spark-3 \  
  --source lakefs://example-repo@main  
# output:  
# created branch 'testing-spark-3.0', pointing to commit ID: 'd1e9adc71c10a'
```

```
CREATE TABLE master.request_logs (  
  request_time timestamp,  
  url varchar,  
  ip varchar,  
  user_agent varchar  
)  
WITH (  
  format = 'TEXTFILE',  
  external_location = 's3a://example/main/data/logs/'  
);
```




# Continuous Data Integration

- Ingest new data safely
- Enforce best practices
- **Metadata validation** - prevent breaking changes from entering your production data environment



# Enforce Best Practices





# CI with PyArrow

```
import lakefs
import pyarrow.parquet as pq

# Setup a PyArrow FileSystem that we can use to query data in the source ref
fs = lakefs.get_filesystem(repo, source_ref)

for change in lakefs.diff(repo, source_ref, target_branch, prefix='public/'):
    if not change.path.endswith('.parquet'):
        continue
    # Read Parquet column metadata
    schema = pq.read_schema(fs.open_input_file(change.path))
    if filter(lambda column: column.name == 'user_id', schema):
        raise ValueError('user_id column not allowed')
```

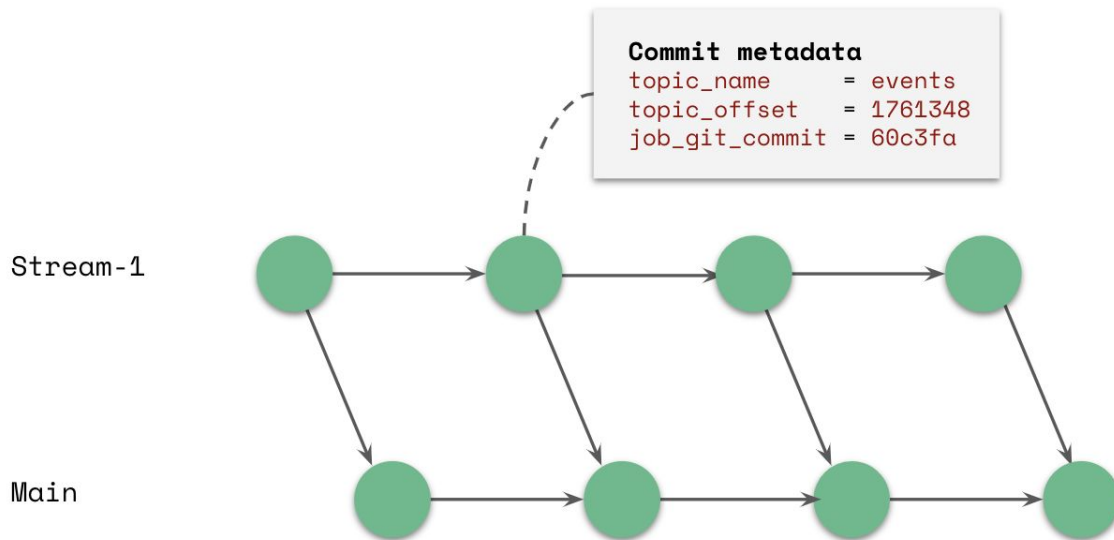


# Continuous Data Deployment

- **Prevent data quality issues** - by testing production data before exposing it to users / consumers
- **Test** - intermediate results in your DAG to avoid cascading quality issues
- **Instantly revert changes to data**



# Streaming Data with Kafka Connect



# Streaming Data with Kafka Connect

```
# Your lakeFS repository
s3.bucket.name=example-repo

# Your lakeFS S3 endpoint and credentials
store.url=https://s3.lakefs.example.com
aws.access.key.id=AKIAIOSFODNN7EXAMPLE
aws.secret.access.key=wJalrXUtnFEMIK7MDENGbPxRfiCYEXAMPLEKEY

# master being the branch we want to write to
topics.dir=example-branch/topics
```

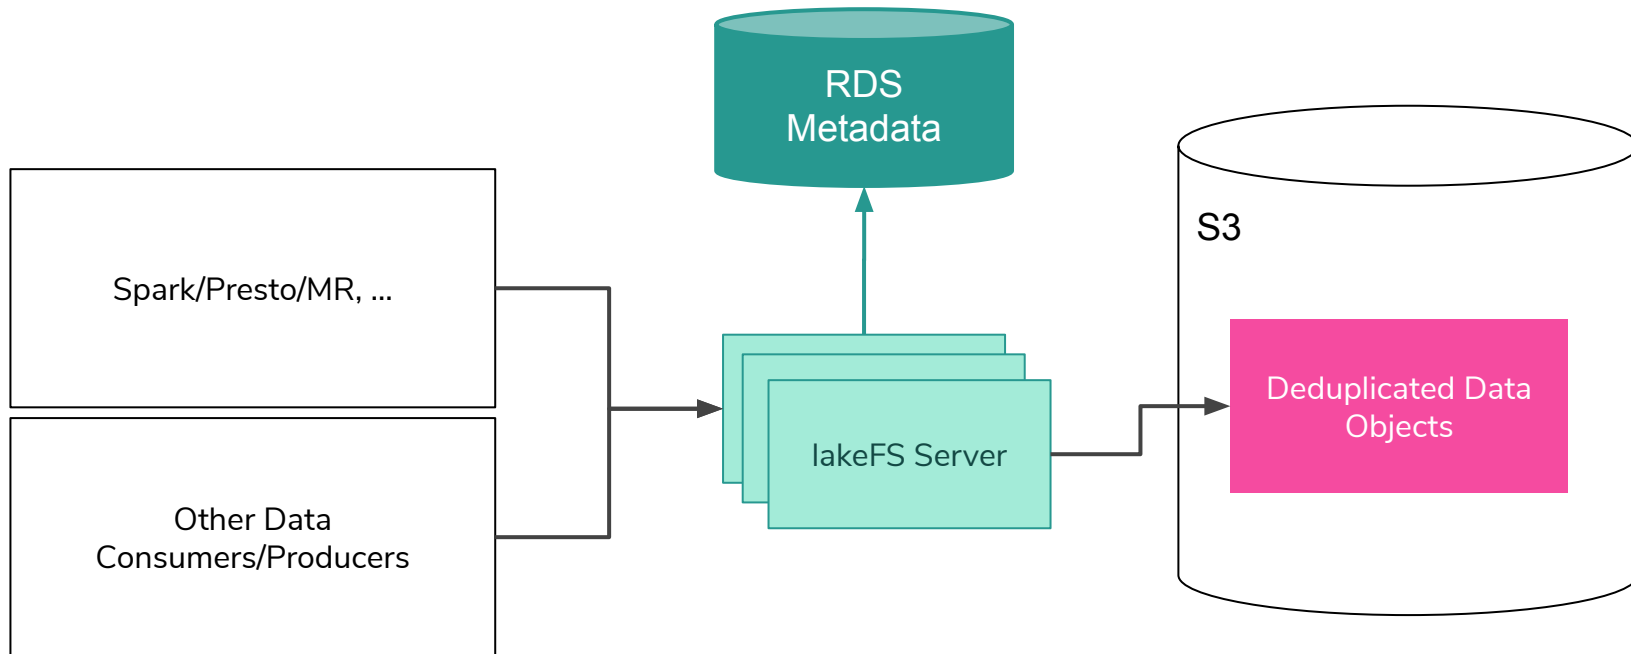
```
Lakectl branch revert lakefs://example-repo@stream-1-branch --commit dd8a60d5ef70809
```



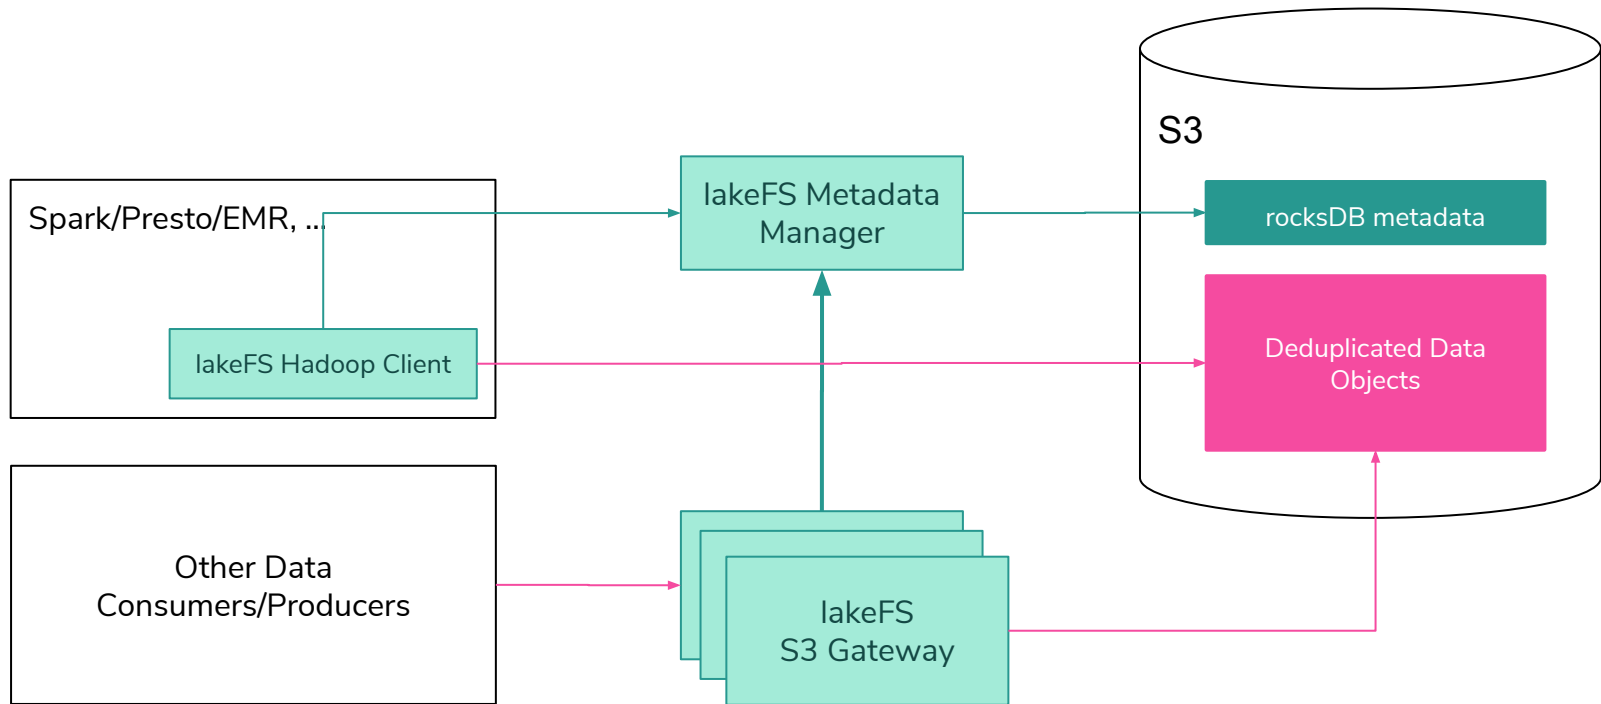
# lakeFS Architecture



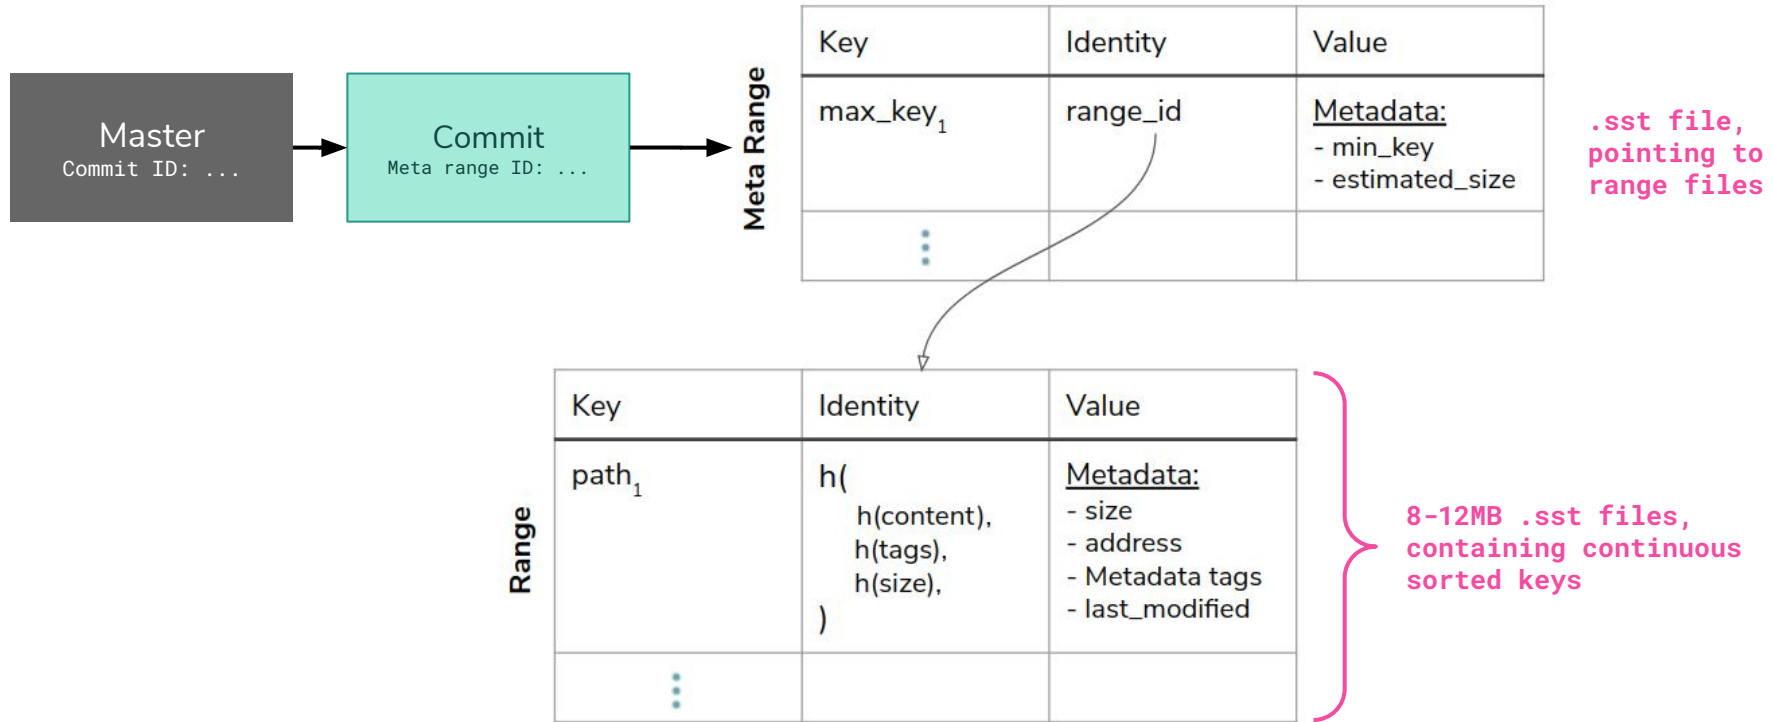
# Old Architecture



# In-Development Architecture



# Future Architecture - Data Format



[Further Reading](#)



# Additional Resources



[Getting started](#)



Check out the [docs](#)



Join the [lakeFS Slack Channel](#)



Contribute and star the [repo](#)



# Thanks!



**lakeFS**



Join the community



GitHub



einat.orr@treeverse.io

# The open core model

